

Kyber1024 | selected kernel throughput | batch = 32,768

Saved GPU timings: batch / cumulative time of each kernel across the CPA pipeline; excludes host API overhead

